

HW4

Bitcoin Miner

Parallel Programming
2025/10/31


Spec

In this assignment, you are going to **parallelize a bitcoin miner with CUDA**. You may use any techniques that you've learned in this class!

Please refer to [the assignment's specification](#) for detailed information about the assignment.

What is Blockchain?

- A blockchain is essentially a sequence of verified **blocks**
 - Each block contains the information of transactions and metadata, including the hash value of its previous block.
 - Due to the property of [cryptographic hash function](#), If something is modified, it will be discovered immediately.



What is Block?

- A block is a data structure containing the information of transactions and metadata. It consists of two parts:
 - Block Header
 - Version number
 - Hash of the previous block
 - Merkle root
 - Timestamp
 - Target Difficulty (encoded in Bits)
 - Nonce
 - Block Body
 - Transaction counter
 - List of transaction records
 - ...

version	02000000
previous block hash (reversed)	17975b97c18ed1f7e255adf297599b55 330edab87803c8170100000000000000
Merkle root (reversed)	8a97295a2747b4f1a0b3948df3990344 c0e19fa6b2b92b3a19c8e6badc141787
timestamp	358b0553
bits	535f0119
nonce	48750833
transaction count	63
coinbase transaction	
transaction	
...	

Block hash
0000000000000000
e067a478024addfe
cdc93628978aa52d
91fabd4292982a50



Job of Bitcoin Miner

The job of a miner is to create and add blocks to the blockchain by following steps:

1. Gather transactions
2. Create the block header
3. Find the proper nonce
4. Broadcast the block
5. Receive rewards


version	02000000
previous block hash (reversed)	17975b97c18ed1f7e255adf297599b55 330edab87803c8170100000000000000
Merkle root (reversed)	8a97295a2747b4f1a0b3948df3990344 c0e19fa6b2b92b3a19c8e6badc141787
timestamp	358b0553
bits	535f0119
nonce	48750833
transaction count	63
coinbase transaction	
transaction	
...	

Block hash
0000000000000000
e067a478024adfe
cdc93628978aa52d
91fabd4292982a50

hash < Target Difficulty

What is Merkle Tree?


A Merkle tree is a binary tree in which every leaf node is the hash of a data block, and every inner node is the hash its children. Merkle tree allows **efficient and secure verification of the contents of a large data structure.**



Node Generation

- Each parent node is generated by **double SHA-256 function**
 1. Each child node is hashed by SHA-256 TWICE.
 2. The results (H_A and H_B) are summed together.
 3. The sum is then hashed AGAIN by SHA-256 TWICE.
- The procedure is repeated again and again, until the root node is reached


$$H_A = \text{SHA256}(\text{SHA256}(TxA)) \quad H_B = \text{SHA256}(\text{SHA256}(TxB))$$



$$H_{AB} = \text{SHA256}(\text{SHA256}(H_A + H_B))$$

Merkle Root

The root is called Merkle root. It represents the entire data set, and any change in the data would result in a different Merkle root.



Credit: <https://www.oreilly.com/library/view/mastering-bitcoin/9781491902639/ch07.html>

Target Difficulty

- Purpose: to maintain a consistent block generation rate and ensure that the blockchain remains secure and stable.
- The Target Difficulty is a value to determine how difficult it is to find out a proper hash for the block header by requiring $\text{hash} < \text{Target Difficulty}$.
- The Target Difficulty conforms to the following rules
 - It is a 256-bit number, encoded by a field called “bits” (in 4 bytes)
 - It has N leading zeros. Currently, N equals 76 (year 2022).
 - It is determined by the Bitcoin network.
 - It is adjusted such that a new block is computed every 10 minutes (on average)
 - You may check current difficulty here: <https://btc.com/stats/diff>

Nonce

- Nonce is a 32-bit value in the block header.
- Miner needs to find a proper nonce such that the hash of the block header is less than the current Target Difficulty.

version	02000000
previous block hash (reversed)	17975b97c18ed1f7e255adf297599b55 330edab87803c8170100000000000000
Merkle root (reversed)	8a97295a2747b4f1a0b3948df3990344 c0e19fa6b2b92b3a19c8e6badc141787
timestamp	358b0553
bits	535f0119
nonce	48750833
transaction count	63
coinbase transaction	
transaction	
...	

Block hash
0000000000000000
e067a478024addfe
cdc93628978aa52d
91fabd4292982a50

Block header hash < Target Difficulty

Goal of the Assignment

- Find a proper nonce for a given block
 - The hash of the block header has to be less than the given Target Difficulty.
 - The time you spent on finding the hash value has to be accelerated.
 - Use any techniques that you've learned in this class!

Sequential Code

1. Read input.
2. Calculate the merkle root from transactions.
3. Decode the Target Difficulty.
4. Find a proper nonce, such that the hash value of the block header satisfies the requirement.

```
for nonce = 0X00000000 to 0Xffffffff  
    calculate a hash value  
    if the hash value < Target Difficulty  
        done
```

Sequential Code - Header Extraction

- Step 1: Read input
- There are several fields in the block header that you have to extract.

```
// ***** read data *****
char version[9];
char prevhash[65];
char ntime[9];
char nbits[9];
int tx;
char *raw_merkle_branch;
char **merkle_branch;

getline(version, 9, fin);
getline(prevhash, 65, fin);
getline(ntime, 9, fin);
getline(nbites, 9, fin);
fscanf(fin, "%d\n", &tx);
printf("start hashing")

raw_merkle_branch = new char [tx * 65];
merkle_branch = new char *[tx];
for(int i=0;i<tx;++i) {
    merkle_branch[i] = raw_merkle_branch + i * 65;
    getline(merkle_branch[i], 65, fin);
    merkle_branch[i][64] = '\0';
}
```

Sequential Code - Header Fields

- A representative header fields of a block is depicted below
 - Please have a comparison with pages 4.
 - These information has to be extracted first before we carrying out the calculation.
- `case01.in`

4	Number of blocks
20000000	Version
00000000000000000003348540cbfc68b70825e7abcd5a83a48a5f87fa7f1aace	Previous block hash
5ac22f8b	Timestamp
17502ab7	Bits (packed difficulty)
2094	Number of transactions
c6574adb277efbf972658ab78b1277707a967076cfcc90d6af800cd8a915396d	Transactions
c01c33240ba97fb6db2b98fbaf7e4211fe3b59585372994bdac1b96b1c9be0d3	
abce7b2b30ff62b4998864db6a1ea77db8eb33d3f6abe3f981e0d2802c999042	
fd463aca59df1302c8f08e8070b56bd64ebe0cf35ba68cd5b39d208a2ff50053	
4321318516cb6630e3b3caa29bd49227168a69d170e1110c220b3c30d15f913c	
f1f452e09dad935a67499a24e388d61fd7f0f38c97a2a92c73b2fd7019a85578	
dd24a8e3f419006cd2f7cb3336b605fa6c2898accecac3d0a844c2ae8f5f53d9	

Sequential Code - Merkle Root

- Step 2: Calculate the merkle root from transactions
 - The merkle root is calculated by a function provided by the TA
 - You are encouraged to take a look of its implementation

```
// ***** calculate merkle root *****

unsigned char merkle_root[32];
calc_merkle_root(merkle_root, tx, merkle_branch);

printf("merkle root(little): ");
print_hex(merkle_root, 32);
printf("\n");

printf("merkle root(big):      ");
print_hex_inverse(merkle_root, 32);
printf("\n");
```

Sequential Code - Target Difficulty

- Step 3: Decode the Target Difficulty
 - It is decoded from a number called nbBits.
 - The decode algorithm is already implemented by us (shown below).
 - If you are interested in that algorithm, please take a look of this [website](#).

```
// ***** calculate target value *****
// calculate target value from encoded difficulty
which is encoded on "nbits"
unsigned int exp = block.nbits >> 24;
unsigned int mant = block.nbits & 0xffffffff;
unsigned char target_hex[32] = {};

unsigned int shift = 8 * (exp - 3);
unsigned int sb = shift / 8;
unsigned int rb = shift % 8;
```

Sequential Code -Nonce

- Step 4 : Find a proper nonce, such that the hash value of the block header satisfies the requirement
 - Try if you can figure out a way to parallelize and accelerate the code.

```
for(block.nonce=0x00000000; block.nonce<=0xffffffff;++block.nonce) {
    //sha256d
    double_sha256(&sha256_ctx, (unsigned char*)&block, sizeof(block));
    if(block.nonce % 1000000 == 0) {
        printf("hash #%-10lu (big): ", block.nonce);
        print_hex_inverse(sha256_ctx.b, 32);
        printf("\n");
    }
    if(little_endian_bit_comparison(sha256_ctx.b, target_hex, 32) < 0) { // sha256_ctx < target_hex
        printf("Found Solution!!\n");
        printf("hash #%-10lu (big): ", block.nonce);
        print_hex_inverse(sha256_ctx.b, 32);
        printf("\n\n");
        break;
    }
}
```

Sequential Code - SHA-256

- SHA-256 implementation for the sequential version can be found at:
 - sha256.h
 - sha256.cu
- If you are interested in the implementation details, please take a look at the following website
 - <https://en.wikipedia.org/wiki/SHA-2>

Test cases

- We provide several test cases for you to test your program
- The execution time of the sequential code are given below:
 - case00.in : about 1 hr 15 min
 - case01.in : about 1 hr 15 min
 - case02.in : about 30 min
 - case03.in : about 11 min
- Please keep in mind that your goal is to parallelize the mining algorithm and accelerating it.
- **It has to be faster than the sequential version.**

Grading

- Correctness (50%)
 - Please make sure to use GPU and verify your answer.
- Performance (20%)
 - We will grade the performance of your assignments against the other students in the class.
- Report (30%)
 - Please describe, in detail, how you parallelize your codes as well as your optimization methodology.
- Advanced CUDA skills
 - You are welcome to use streaming, page-lock memory, asynchronous memory copy, or any other advanced skills.
- Others
 - You will get credits if you are able to optimize the other parts of the source codes. Please justify your solutions and provide a detailed comparison in your report.

Grading - Report

- In your report, please include the following parts:
 - Describe the design/overview of your implementation briefly.
 - The parallelization and optimization techniques you used in your solution
 - Experiments of various combinations of the number of blocks & threads (at least 8 combinations) and plot them with the figures
 - Describe the details if you use advanced CUDA skills
 - (Optional) Any suggestions or feedback for the homework is welcome

Submission

- Please zip the following files to a single archived file named <StudentID>.tar:
 - hw4.cu
 - sha256.cu
 - sha256.h
 - report.pdf
 - Makefile (optional)
- Please make sure your Makefile works properly (Your Makefile should build executable binary hw4) and your program can run before submitting your assignment 4.

Reminder

- Because we are doing hash, there may be a situation multiple nonce can satisfies the requirement hash value < Target Difficulty.
- We accept all nonce satisfies the requirement, so don't worry if your nonce isn't same as our provided solutions.
- We ensure that all the testcases have more than one solution.
- Please refer to [the assignment's specification](#) for detailed information about the assignment.

Deadline

- The deadline of assignment 4 is **11/14 (Fri.), 23:59.**
- Everyone is welcomed to ask questions on NTU cool.